

# COMPSCI 589 Final Project - Spring 2026

Due **May 8**, 11:55 pm Eastern Time

## Introduction

- Each group member was responsible for evaluating the performance of at least one unique algorithm implemented during the semester on each dataset
- The algorithms tested in this project were chosen from the pool of:  $k$ -NN, Decision Trees, multinomial Naive Bayes, Random Forests, and Neural Networks
- All work was done in the shared github repository [https://github.com/19851xw/589\\_final\\_project](https://github.com/19851xw/589_final_project)

- 
- **Katherine** implemented TODO
  - **Tiffany** implemented TODO
  - **Xing-Wei** implemented Decision Trees

# Experiments and Analyses

## 1 The Hand-Written Digits Recognition Dataset

### 1.1 Neural Networks Algorithm

TODO

---

### 1.2 Algorithm 2

TODO

---

### 1.3 Algorithm 3

TODO

## **2   The Oxford Parkinson's Disease Detection Dataset**

### **2.1   Neural Networks Algorithm**

TODO

---

### **2.2   Algorithm 2**

TODO

---

### **2.3   Algorithm 3**

TODO

## 3 The Rice Grain Dataset

### 3.1 Neural Networks Algorithm

TODO

---

### 3.2 Algorithm 2

TODO

---

### 3.3 Decision Trees Algorithm

We chose to test decision trees on the rice dataset because it has a moderate number of features (7) and a relatively large number of samples (3810), which allows us to explore how different tree depths and minimum split sizes affect performance.

Tuning for dataset with 3810 rows and 7 features.

Testing Depths: [2, 3, 5, 7, 9, 11]

Testing Splits: [190, 381, 571, 762]

| Depth | Split | Test Acc | Test F1 |
|-------|-------|----------|---------|
| 2     | 190   | 0.9215   | 0.9192  |
| 2     | 381   | 0.9228   | 0.9205  |
| 2     | 571   | 0.9226   | 0.9202  |
| 2     | 762   | 0.9239   | 0.9215  |
| 3     | 190   | 0.9234   | 0.9210  |
| 3     | 381   | 0.9213   | 0.9189  |
| 3     | 571   | 0.9202   | 0.9179  |
| 3     | 762   | 0.9239   | 0.9215  |
| 5     | 190   | 0.9213   | 0.9189  |
| 5     | 381   | 0.9234   | 0.9210  |
| 5     | 571   | 0.9226   | 0.9202  |
| 5     | 762   | 0.9239   | 0.9215  |
| 7     | 190   | 0.9223   | 0.9199  |
| 7     | 381   | 0.9226   | 0.9203  |
| 7     | 571   | 0.9215   | 0.9192  |
| 7     | 762   | 0.9239   | 0.9215  |
| 9     | 190   | 0.9234   | 0.9210  |
| 9     | 381   | 0.9220   | 0.9197  |
| 9     | 571   | 0.9239   | 0.9215  |
| 9     | 762   | 0.9239   | 0.9215  |
| 11    | 190   | 0.9223   | 0.9199  |
| 11    | 381   | 0.9220   | 0.9197  |
| 11    | 571   | 0.9218   | 0.9194  |
| 11    | 762   | 0.9239   | 0.9215  |

=====

BEST CONFIGURATION:

Max Depth: 11

Min Split: 762

F1 Score: 0.9215

=====

Final Results:

Training Accuracy: 0.9239

Training F1 Score: 0.9216

Testing Accuracy: 0.9239

Testing F1 Score: 0.9215




Figure 1: Rice Accuracy Distribution




Figure 2: Rice F1 Score Distribution

For both models, the training scores are highly concentrated around 0.92, while the testing scores are more spread out, with a mean around 0.923 for accuracy and 0.922 for F1 score. The training performance suggests that the models have high stability and low bias, however, the testing performance indicates that there may be some sensitivity to the specific training data used in each fold. Since the mean training scores are very close to the mean testing scores, it suggests that the models are robust and are not overfitting. Decision trees can easily split numerical features to find precise boundaries, so they are well-suited for the rice dataset, which consists of numerical attributes.

## 4 The Credit Approval Dataset

### 4.1 Neural Networks Algorithm

TODO

---

### 4.2 Algorithm 2

TODO

---

### 4.3 Algorithm 3

TODO

## Extra Credit

We did the following extra credit tasks. Describe what you did as part of that task and present the corresponding results. Any source code you develop should be included in your Gradescope submission.

### [QE.1] [Extra Credit Question #1] (10 Points)

You may earn up to 10 extra points if you analyze all four datasets using an additional algorithm. In other words, you would be evaluating more than two algorithms on each of the datasets discussed in Section 1.

### [QE.2] [Extra Credit Question #2] (10 Points)

You may earn up to 10 extra points if you select a new challenging dataset and evaluate the performance of different algorithms on it. A challenging dataset should necessarily include both categorical and numerical attributes and have more than two classes; or it should be a dataset where you have to classify images.

### [QE.3] [Extra Credit Question #3] (15 Points)

You may earn up to 15 extra points if you combine different algorithms and construct an ensemble. For instance, you could create an ensemble composed of three neural networks (each with a different architecture) and a random forest. The ensemble would be trained as described in class: each algorithm would be trained on its own bootstrap dataset, and so on. The final predictions would then be determined by majority voting. Evaluate your ensemble algorithm on all four datasets discussed in Section 1.

### [QE.3] [Extra Credit Question #3] (15 Points)

You may earn up to 15 extra points if you implement a new type of algorithm—a non-trivial variant of one of the algorithms studied in class. If you choose to do so, you should evaluate its performance on all four datasets.